

ENABLED TALENT — INTERNAL

CTO Approval Brief

Web Platform Modernization

Prepared for: CTO

Date: March 31, 2026

Status:

FOR APPROVAL

WHAT THIS DOCUMENT IS

This document requests approval to evolve the current website codebase from a marketing-heavy frontend into a fullstack web platform layer that supports:

- Structured content management
- Internal careers and applicant tracking workflows
- Secure admin operations
- Future tenant-aware knowledge experiences
- Future company-specific chatbot workflows

This is not a replacement for the core product platform evolution plan. It is the web experience and content operations layer that will sit beside it.

WHAT WE ARE APPROVING

We are approving a phased modernization of this repository around:

- Next.js 16 App Router
- Supabase for auth, Postgres, storage, and row-level access control
- Tiptap OSS core for Medium-style editorial authoring
- OpenAI Responses API for later retrieval-grounded chatbot workflows

WHY THIS WORK IS NEEDED

The current repository is increasingly carrying product-like responsibilities, but it still behaves like a static or semi-static marketing site.

Current pain points:

- Content management is limited and inconsistent
- Admin auth is not production-grade
- Blog and careers workflows are not modeled as durable platform features
- Data and content operations are fragmented
- The codebase has no clean path to knowledge ingestion or grounded chat

THE STRATEGIC INTENT

This work aligns with the broader company direction of moving away from a simple marketplace presentation layer and toward a configurable infrastructure experience.

For this repository, that means:

- Making the web layer operationally real
- Treating content and careers as managed system data
- Preparing for tenant-aware experiences without overloading this phase
- Separating platform capability from one-off page implementations

THE ARCHITECTURAL SHIFT

TODAY	TOMORROW
Marketing-heavy frontend	Fullstack web platform layer
Ad hoc admin workflows	Secure authenticated admin workflows
Limited content editing	Structured editorial publishing
Form submissions as isolated utilities	Unified operational data model
No document knowledge layer	Knowledge-ready document pipeline
Generic website support potential	Future tenant-aware assistant support

WHAT WILL BE BUILT

PHASE 1

Platform Foundation

- Supabase auth and role-based admin access
- Core relational schema for content, career postings, applications, inquiries, and media
- Storage buckets for blog images, resumes, and assets
- Replacement of fragile admin patterns with server-side controlled workflows
- Role-based access model (see table below)

PHASE 2

Content and Careers Operations

- Tiptap-based article editor with draft and publish workflow
- Blog CMS supporting two content types:
 - Articles — full posts authored internally using the Tiptap editor
 - External links — curated third-party content (press mentions, partner articles) that appear in the blog grid but open the original source in a new tab
- Careers CMS for internal career postings
- Inquiry and submission management
- Redesigned public blog and careers experiences on live data
- Guest writers can create drafts only; an editor or super admin must review and publish

PHASE 3

Knowledge Readiness

- Company document upload
- Parsing, chunking, and embedding pipeline
- Tenant-aware knowledge document model
- Retrieval-ready storage for future assistants

PHASE 4

Tenant-Aware Assistant Layer

- Grounded chatbot using approved documents
- Company-specific retrieval scopes
- Citation-first answer strategy
- Escalation to human support when confidence is low

Phase 1 — Role-Based Access Model

ROLE	BLOG	CAREERS	ADMIN PANEL	USER MANAGEMENT
Super Admin	Full CRUD, publish and unpublish	Full CRUD	Full access	Can invite and remove all roles

ROLE	BLOG	CAREERS	ADMIN PANEL	USER MANAGEMENT
Editor	Full CRUD, publish	View only	Blog section only	No
Guest Writer	Draft only, cannot publish	No access	Own drafts only	No

Super Admin invites users via email. Role assignment is enforced at the database level through Supabase Row-Level Security policies. Editors carry a team tag (e.g. marketing, engineering) for content attribution, but share the same permission set.

WHAT THIS WORK EXPLICITLY DOES NOT DO

This proposal does not attempt to replace the main product engineering roadmap.

Out of scope:

- Rewriting the core product backend
- Redefining the company-wide auth strategy for the main platform
- Rebuilding all tenant portal flows in this repository
- Implementing full institutional white-label delivery in this phase
- Shipping a company-specific chatbot before the knowledge pipeline is stable

DESIGN PRINCIPLES

The implementation should follow these principles:

- Tenant awareness without unnecessary early complexity
- Secure defaults for auth, storage, and content access
- One reusable operational model instead of many page-level exceptions
- Editorial quality equal to modern B2B publishing tools
- Phased delivery so value is created before advanced AI layers are added

EXPECTED OUTCOMES

At the end of the approved scope, this repository should provide:

- A real admin layer instead of a pseudo-admin surface

- Structured publishing and content operations
- Durable careers and inquiry workflows
- A stronger public-facing product and content surface
- A clean base for later tenant-scoped knowledge and assistant features

KEY RISKS

Risk 1 — Scope Creep

There is a risk that this work tries to solve both the website modernization and the core platform transformation at the same time.

Mitigation:

- Keep this proposal scoped to the web experience layer
- Treat tenant-aware knowledge and chat as later phases
- Keep the main product platform plan separate

Risk 2 — Auth and Data Fragmentation

If the web layer evolves with no clean operational model, the result will be more fragmented systems rather than fewer.

Mitigation:

- Centralize auth, storage, and content workflows
- Use one shared schema and role model for this repository

Risk 3 — AI Before Data Readiness

A chatbot added before content quality, permissions, and document grounding are stable will create poor answers and operational risk.

Mitigation:

- Phase AI after the content and knowledge base are stable
- Use retrieval grounding rather than generic responses

APPROVAL REQUEST

1 The Scope Boundary

This repository will be modernized as a web platform layer, not as a total replacement for the main product platform.

2 The Stack

Approve Next.js + Supabase + Tiptap + OpenAI for this repository's modernization path.

3 The Phased Sequence

Approve foundation first, then CMS and workflows, then knowledge ingestion, then chatbot.

4 The Operating Principle

Treat this codebase as a reusable web experience platform, not only as a set of marketing pages.

SUCCESS CRITERIA

This proposal should be considered successful if it delivers:

- Secure admin access
- Live content and careers management
- A publishable article workflow
- Operational inquiry handling
- A knowledge-ready architecture for future assistant use

RECOMMENDED DECISION

Approve the modernization of this repository as a phased web platform program, with immediate focus on auth, data model, CMS, and operations, while deferring tenant-specific assistant behavior until the knowledge layer is ready.

APPROVED BY

PREPARED BY

CTO · Date: _____

Engineering · Date: March 31, 2026

CONFIDENTIAL — ENABLED TALENT INTERNAL USE ONLY